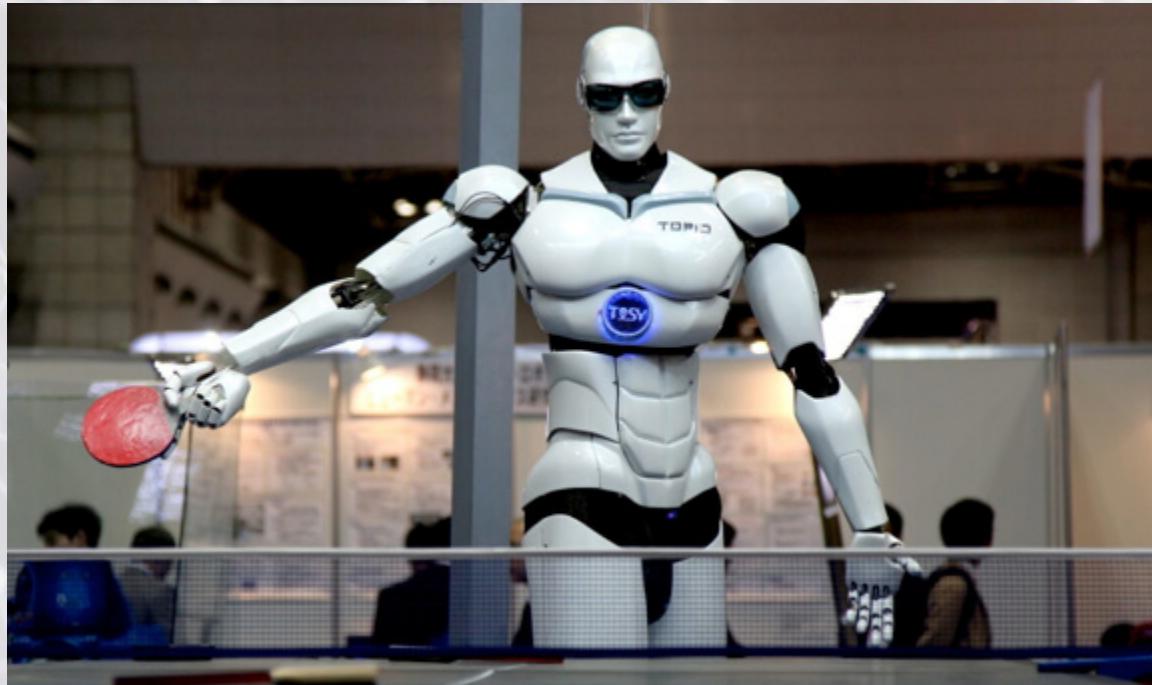
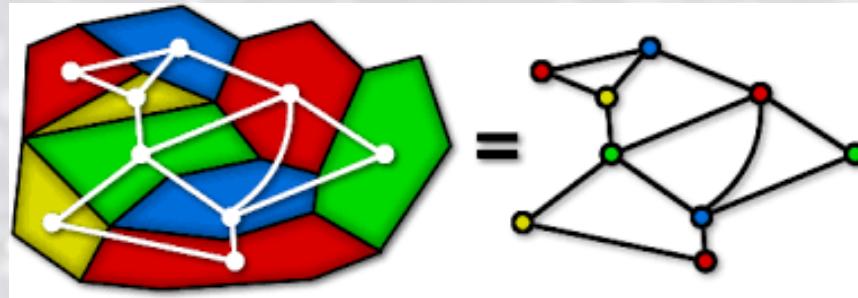


Artificial Intelligence

Chapter 6: Constraint Satisfaction Problems



Digression: The Four Color Theorem

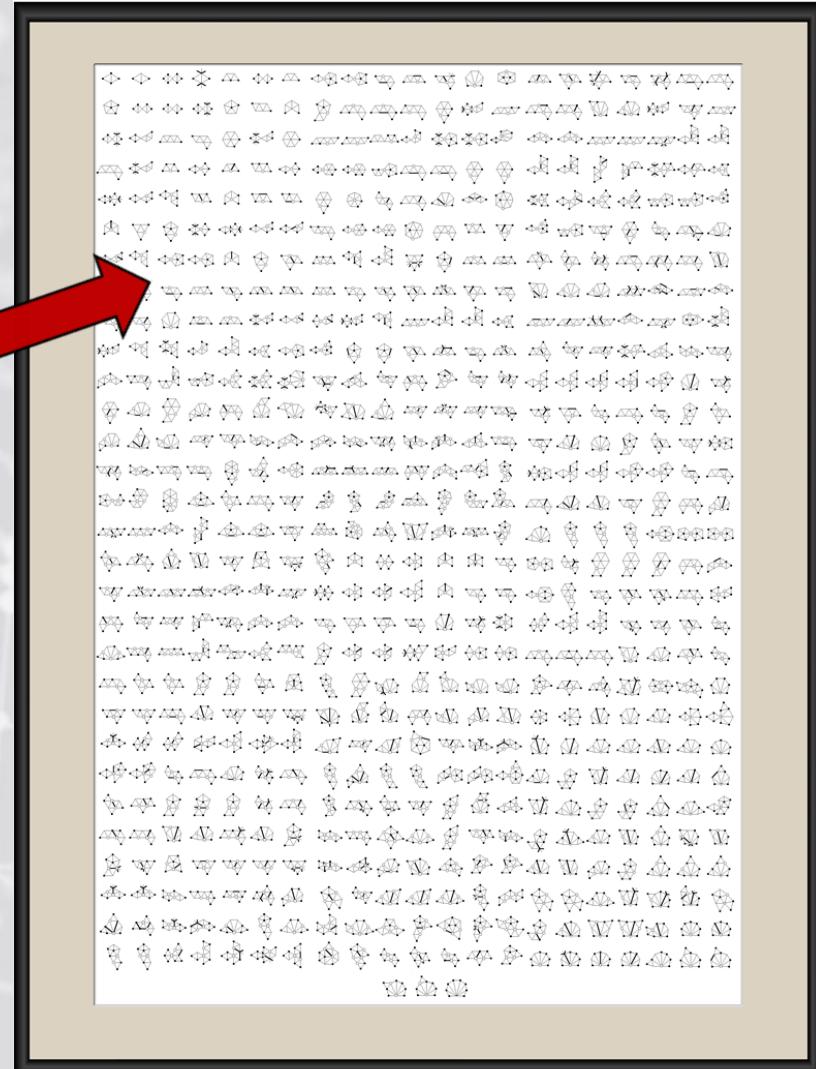


- One of the most famous results in the history of mathematics.
- It was first conjectured in 1852, but only finally proven in 1976. Notably it was the first math proof to rely crucially on computers (for a large set of configuration/case checks) – and for this reason was considered controversial.

Digression: The Four Color Theorem

- Why was the proof so difficult?

Because the best-known technique relied on (originally) 1936 unavoidable configurations -- like these.

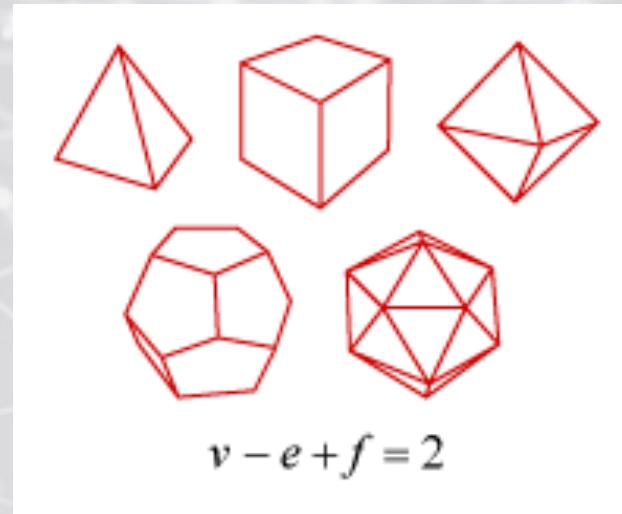


Digression: The Four Color Theorem

- While we can't prove the FCT in lecture (!), we can prove its baby brother, *the 6 color theorem*.
- Let's do this...
- (1) Prove Euler's Polyhedron Formula.

Digression: The Four Color Theorem

- (1) Prove Euler's Polyhedron Formula. (we use induction on $E(G)$).



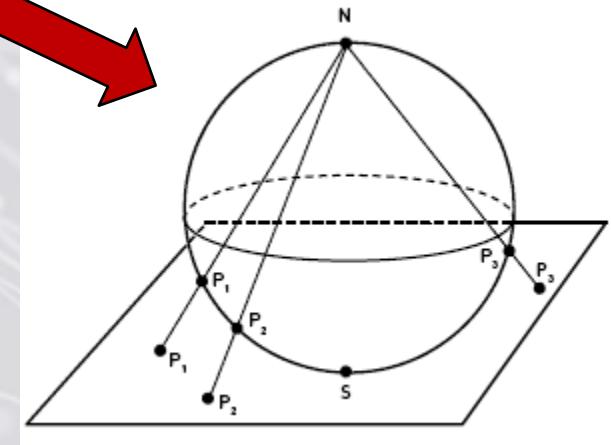
Digression: The Four Color Theorem

- (1) Prove Euler's Polyhedron Formula.
- (2) Convince yourself that embedding on the sphere is equivalent to planarity.

Digression: The Four Color Theorem

- (1) Prove Euler's Polyhedron Formula.
- (2) Convince yourself that embedding on the sphere is equivalent to planarity.

This is called a
“stereographic
projection”



Digression: The Four Color Theorem

- (1) Prove Euler's Polyhedron Formula.
- (2) Convince yourself that embedding on the sphere is equivalent to planarity.
- (3) Prove that if G is planar (with $n \geq 3$), then $|E(G)| \leq 3n - 6$.

Digression: The Four Color Theorem

- (1) Prove Euler's Polyhedron Formula.
- (2) Convince yourself that embedding on the sphere is equivalent to planarity.
- (3) Prove that if G is planar (with $n \geq 3$), then $|E(G)| \leq 3n - 6$.
- (4) Claim: Every planar graph contains a vertex of degree 5 or less.

Digression: The Four Color Theorem

- (1) Prove Euler's Polyhedron Formula.
- (2) Convince yourself that embedding on the sphere is equivalent to planarity.
- (3) Prove that if G is planar (with $n \geq 3$), then $|E(G)| \leq 3n - 6$.
- (4) Claim: Every planar graph contains a vertex of degree 5 or less. (Hint: use the fact that the sum of the degrees of vertices in a graph equals twice the number of edges).

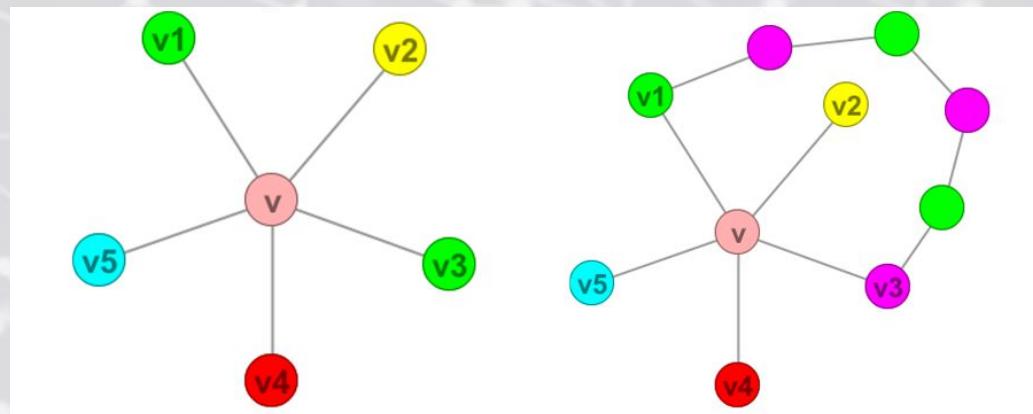
Digression: The Four Color Theorem

- (1) Prove Euler's Polyhedron Formula.
- (2) Convince yourself that embedding on the sphere is equivalent to planarity.
- (3) Prove that if G is planar (with $n \geq 3$), then $|E(G)| \leq 3n - 6$.
- (4) Claim: Every planar graph contains a vertex of degree 5 or less.
- Now put it all together and we have the 6 Color Theorem!



Digression: The Four Color Theorem

- So how about a Five Color Theorem?
- Actually, it's not *that* bad. We use the Six Color Theorem plus an additional clever argument utilizing a structure called a “Kempe Chain” (1890); the result follows by contradiction.



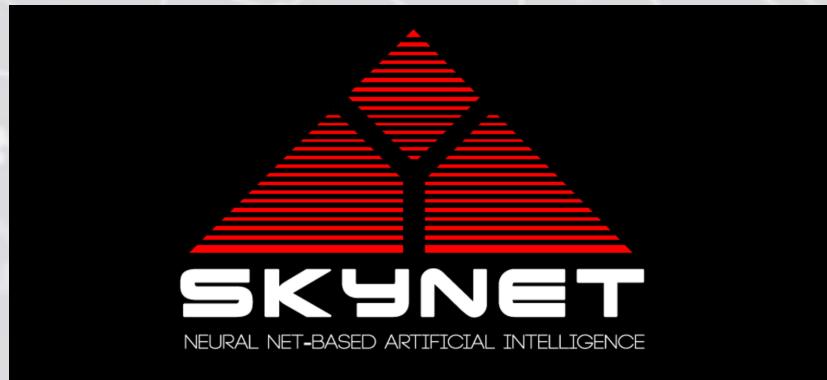
Digression: The Four Color Theorem

- Unfortunately, no one to date has found a simple way to reduce Kempe Chains and similar structures to efficiently solve 4CT (aside from exhaustive case-checking).
- Nevertheless, the 4CT and its proof serve evidence that humankind and machines can work together productively and harmoniously!

Digression: The Four Color Theorem

- Unfortunately, no one to date has found a simple way to reduce Kempe Chains and similar structures to efficiently solve 4CT (aside from exhaustive case-checking).
- Nevertheless, the 4CT and its proof serve evidence that humankind and machines can work together productively and harmoniously!

*This message sponsored by generous donations from your friends at:



Constraint satisfaction problems (CSPs)

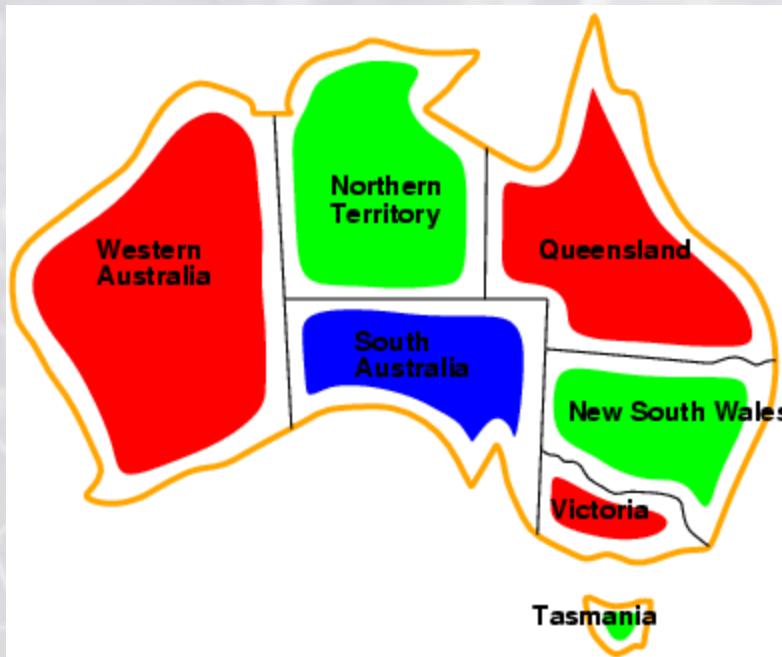
- Standard search problem: **state** is a "black box" – any data structure that supports successor function and goal test
- CSP:
 - **state** is defined by **variables** X_i with **values** from **domain** D_i
 - **goal test** is a set of **constraints** specifying allowable combinations of values for subsets of variables
- » Allows useful **general-purpose** algorithms with more power than standard search algorithms.

Example: Map-Coloring



- Variables WA, NT, Q, NSW, V, SA, T
- Domains $D_i = \{\text{red}, \text{green}, \text{blue}\}$
- Constraints: adjacent regions must have different colors
 - » e.g., $WA \neq NT$, or $(WA, NT) \in \{(\text{red}, \text{green}), (\text{red}, \text{blue}), (\text{green}, \text{red}), (\text{green}, \text{blue}), (\text{blue}, \text{red}), (\text{blue}, \text{green})\}$

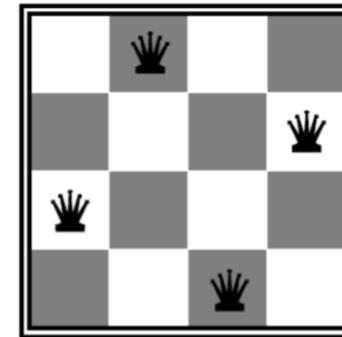
Example: Map-Coloring



- Solutions are **complete** and **consistent** assignments
- e.g., WA = red, NT = green, Q = red, NSW = green, V = red, SA = blue, T = green

Example: N-Queens

- Formulation 1:
 - Variables: X_{ij}
 - Domains: $\{0, 1\}$
 - Constraints



$$\forall i, j, k \quad (X_{ij}, X_{ik}) \in \{(0,0), (0,1), (1,0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{kj}) \in \{(0,0), (0,1), (1,0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k,j+k}) \in \{(0,0), (0,1), (1,0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k,j-k}) \in \{(0,0), (0,1), (1,0)\}$$

$$\sum_{i,j} X_{ij} = N$$

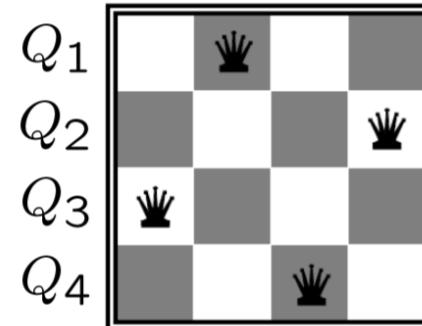
Example: N-Queens

- Formulation 2:
 - Variables: Q_k
 - Domains: $\{1, 2, 3, \dots, N\}$
 - Constraints:

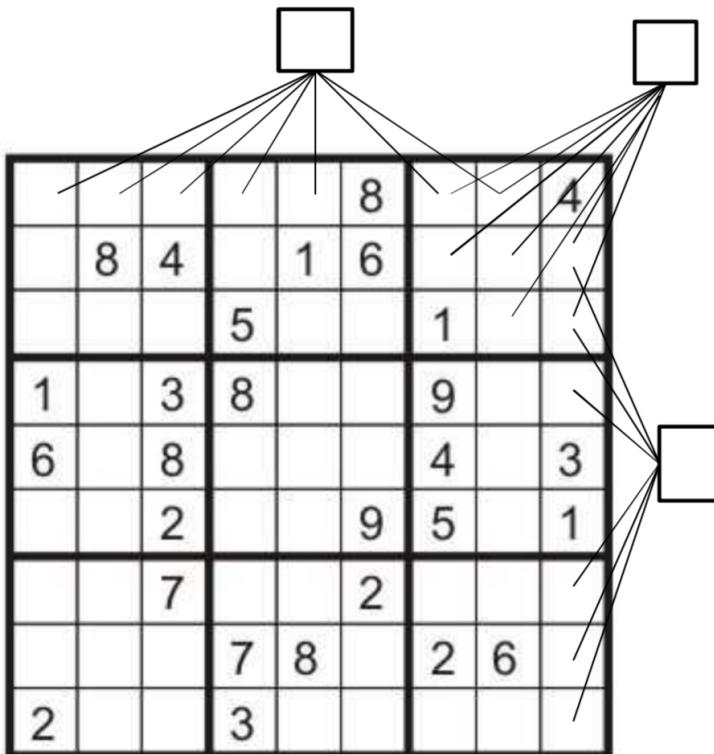
Implicit: $\forall i, j \text{ non-threatening}(Q_i, Q_j)$

-or-

Explicit: $(Q_1, Q_2) \in \{(1, 3), (1, 4), \dots\}$
 \dots



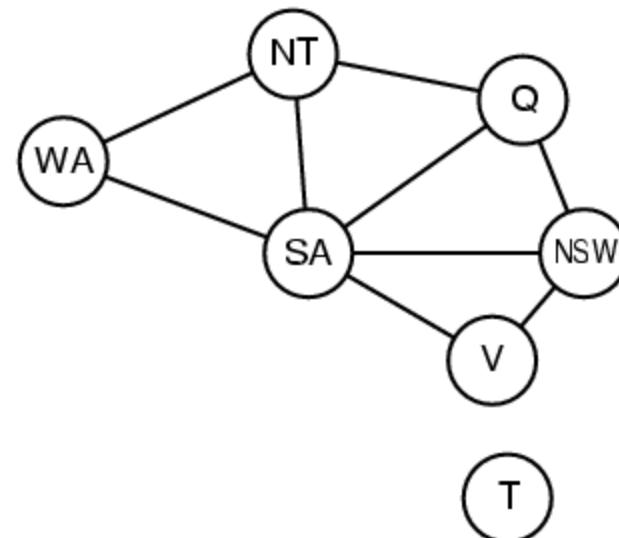
Example: Sudoku



- Variables:
 - Each (open) square
- Domains:
 - $\{1, 2, \dots, 9\}$
- Constraints:
 - 9-way alldiff for each column
 - 9-way alldiff for each row
 - 9-way alldiff for each region
 - (or can have a bunch of pairwise inequality constraints)

Constraint graph

- **Binary CSP:** each constraint relates two variables
- **Constraint graph:** nodes are variables, arcs are constraints



Varieties of CSPs

- Discrete variables
 - finite domains:
 - n variables, domain size $d \rightarrow O(d^n)$ complete assignments
 - e.g., Boolean CSPs, incl. Boolean satisfiability (NP-complete)
 - infinite domains:
 - integers, strings, etc.
 - e.g., job scheduling, variables are start/end days for each job
 - need a constraint language, e.g., $StartJob_1 + 5 \leq StartJob_3$
- Continuous variables
 - e.g., start/end times for Hubble Space Telescope observations
 - linear constraints solvable in polynomial time by LP

Varieties of constraints

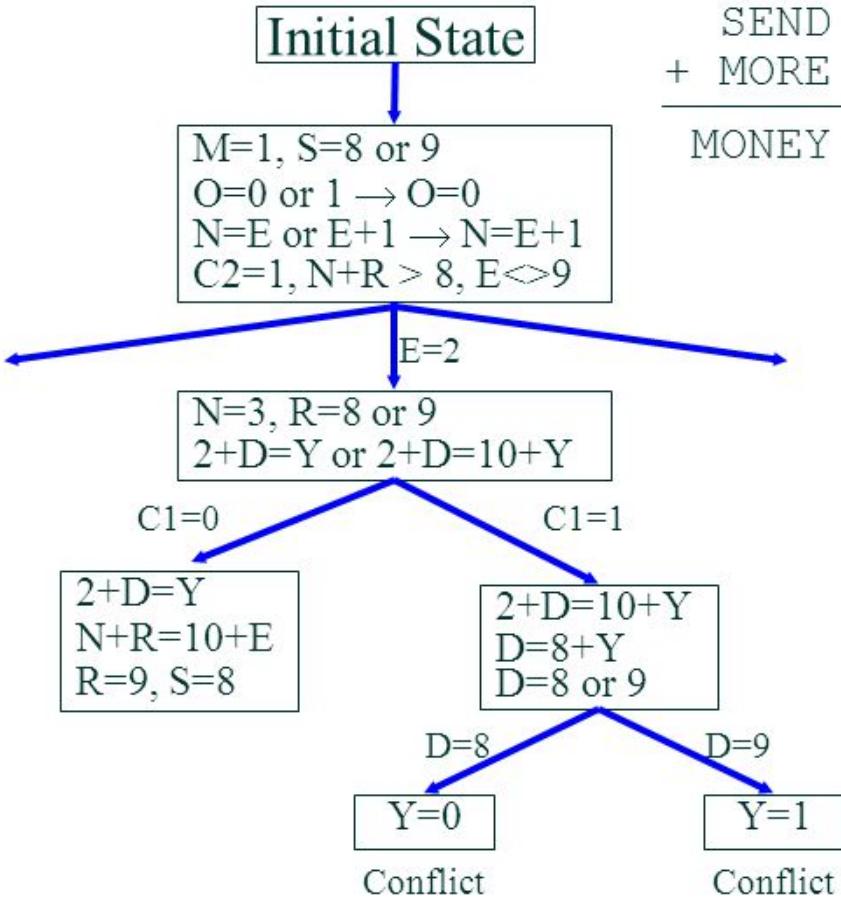
- **Unary** constraints involve a single variable,
 - e.g., SA \neq green
- **Binary** constraints involve pairs of variables,
 - e.g., SA \neq WA
- **Higher-order** constraints involve 3 or more variables,
 - e.g., cryptarithmetic column constraints

Cryptarithmetic Problem

- Try this one:

$$\begin{array}{r} \text{SEND} \\ + \text{MORE} \\ \hline \text{MONEY} \end{array}$$

Cryptarithmetic Problem



- ❑ A useful heuristic can help to select the **best guess** to try first.
- ❑ If there is a letter that participate in many constraints, then it is a good idea to prefer it to a letter that participates in a few.

$$\begin{array}{r}
 \text{SEND} \\
 + \text{MORE} \\
 \hline
 \text{MONEY}
 \end{array}$$

Solution:

$$\begin{array}{r}
 9567 \\
 + 1085 \\
 \hline
 10652
 \end{array}$$

Backtracking search

- Backtracking search is used for a **depth-first search** that chooses values for one variable at a time and backtracks when a variable has no legal values left to assign.
- It repeatedly chooses an unassigned variable, and then tries all values in the domain of that variable in turn, trying to find a solution.
- If an inconsistency is detected, then BACKTRACK returns failure, causing the previous call to try another value.

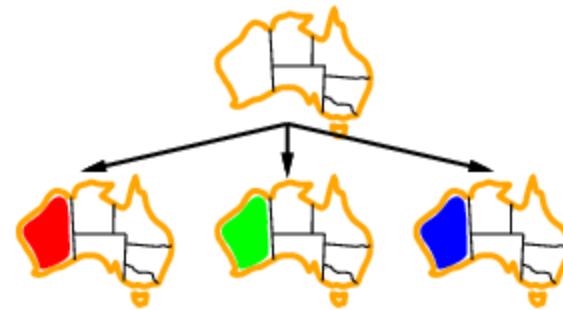
Backtracking Search

- Depth-first search for CSPs with single-variable assignments is called **backtracking** search.
- Variable assignments are **commutative**, i.e.,
[WA = red then NT = green] same as [NT = green then WA = red].
- => Only need to consider assignments to a single variable at each node, so the algorithm keeps only a single representation of a state and alters that representation rather than creating new ones.
 - » Can solve n -queens for $n \approx 25$.

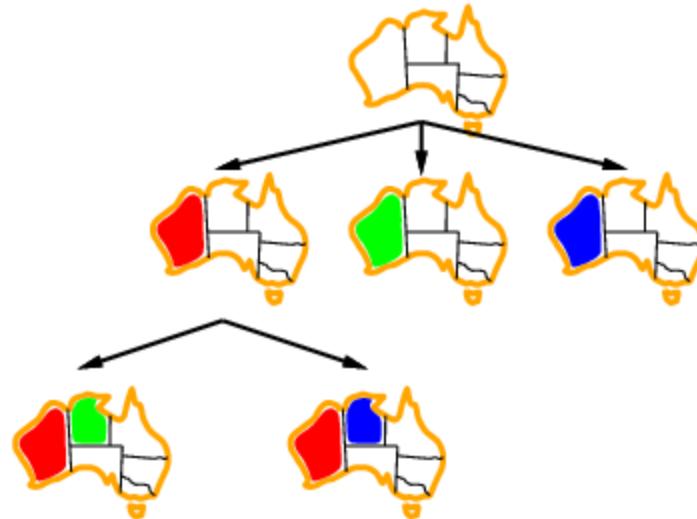
Backtracking example



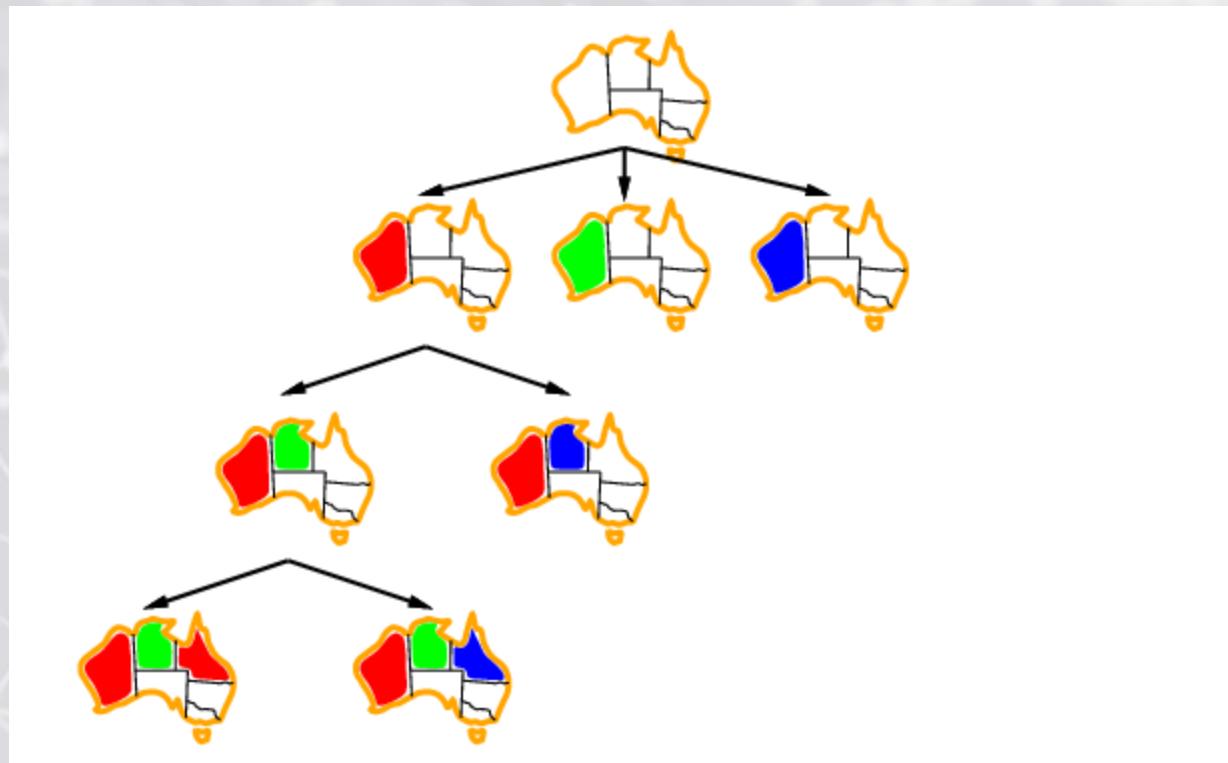
Backtracking example



Backtracking example



Backtracking example



Improving backtracking efficiency

- (3) General-purpose methods can give huge gains in speed:
 - Which variable should be assigned next? (e.g. MRV)
 - In what order should its values be tried? (e.g. inference/forward checking)
 - Can we detect inevitable failure early? (e.g. constraint learning)

Constraint Propagation

- In regular state-space search, an algorithm can do only one thing: *search*.
- In CSP there is a choice: (1) an algorithm can *search* (i.e. choose a new variable assignment from several possibilities); (2) perform a specific type of *inference* called **constraint propagation**.

Constraint Propagation

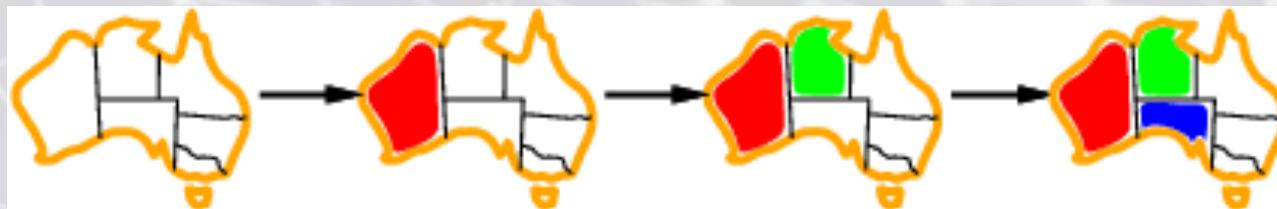
- Constraint propagation uses the constraints to reduce the number of legal values for a variable, which in turn can reduce the legal values for another variable, and so on.
- Constraint propagation may be interleaved with search, or it can be done as a preprocessing step.
- The key idea is: **local consistency**.

Constraint Propagation

- The key idea is: **local consistency** (e.g. node consistency, arc consistency, etc.).
- If we treat each variable as a **node** in the graph, and each binary constraint as an **arc**, then the process of enforcing local consistency in each part of the graph causes inconsistent values to be eliminated throughout the graph.

Most constrained variable

- Most constrained variable:
choose the variable with the fewest legal values



» a.k.a. minimum remaining values (MRV) heuristic

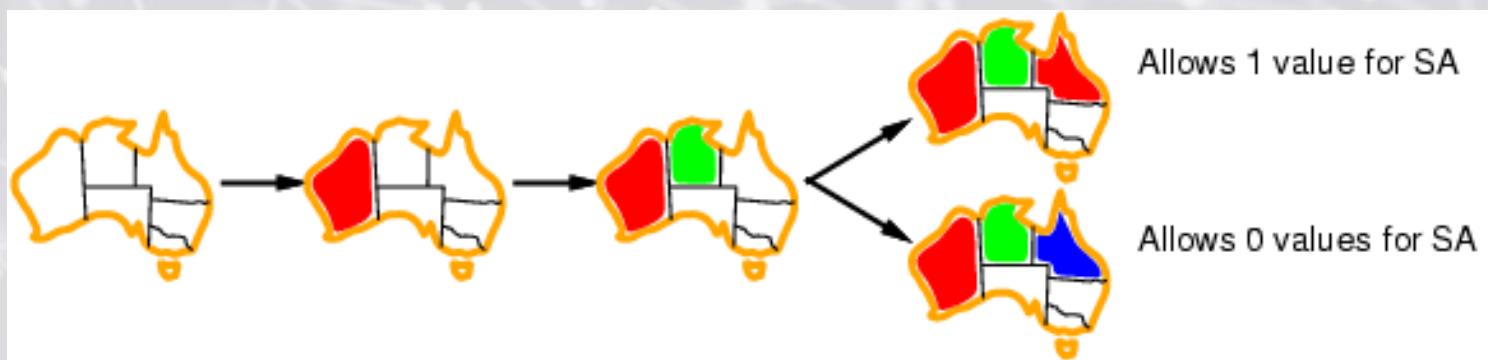
Most constraining variable

- A good idea is to use it as a tie-breaker among most constrained variables
- Most constraining variable:
 - choose the variable with the most constraints on remaining variables



Least constraining value

- Given a variable to assign, choose the least constraining value:
 - the one that rules out the fewest values in the remaining variables



- » Combining these heuristics (MRV + LCV) makes 1000 queens feasible!

Forward checking

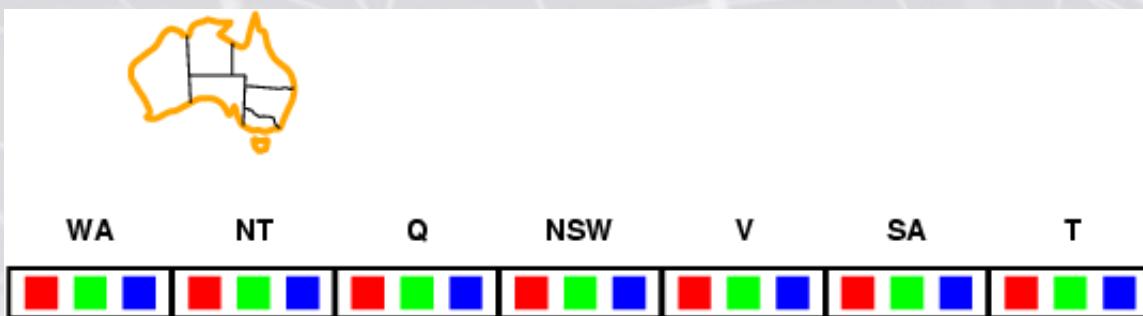
- MRV and LCV and related approaches can infer reduction in the domain variables *before* we begin the search.
- But inference can be even more powerful in the course of a search: every time we make a choice of a value for a variable, we have a brand-new opportunity to infer new domain reductions on neighboring values.

Forward checking

- Forward checking is one of the simplest forms of inference.
- Whenever a variable X is assigned, the forward-checking process establishes **arc consistency** for it: for each unassigned variable Y that is connected to X by a constraint, delete from Y 's domain any value that is inconsistent with the value chosen for X .

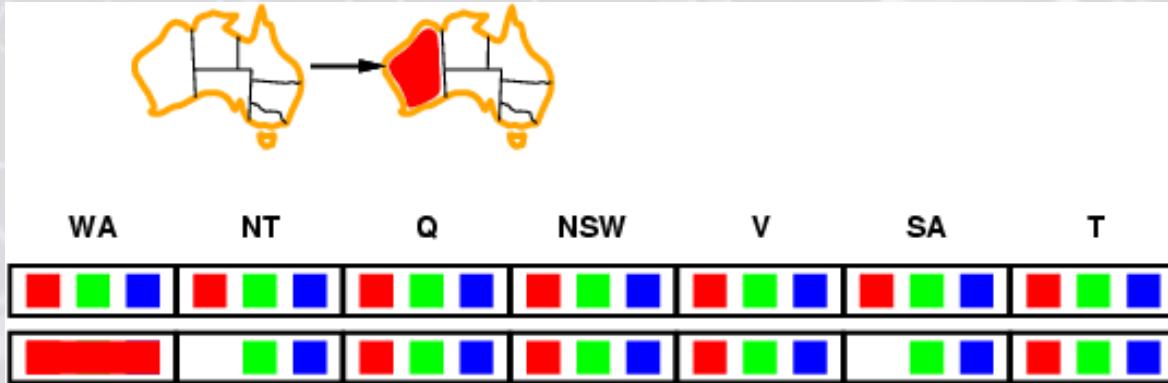
Forward checking

- Idea:
 - Keep track of remaining legal values for unassigned variables.
 - Terminate search when any variable has no legal values.



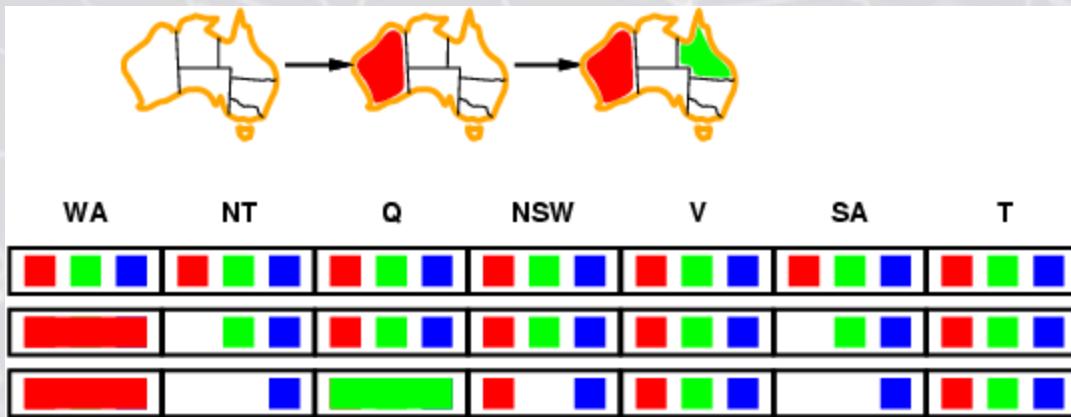
Forward checking

- Idea:
 - Keep track of remaining legal values for unassigned variables.
 - Terminate search when any variable has no legal values.



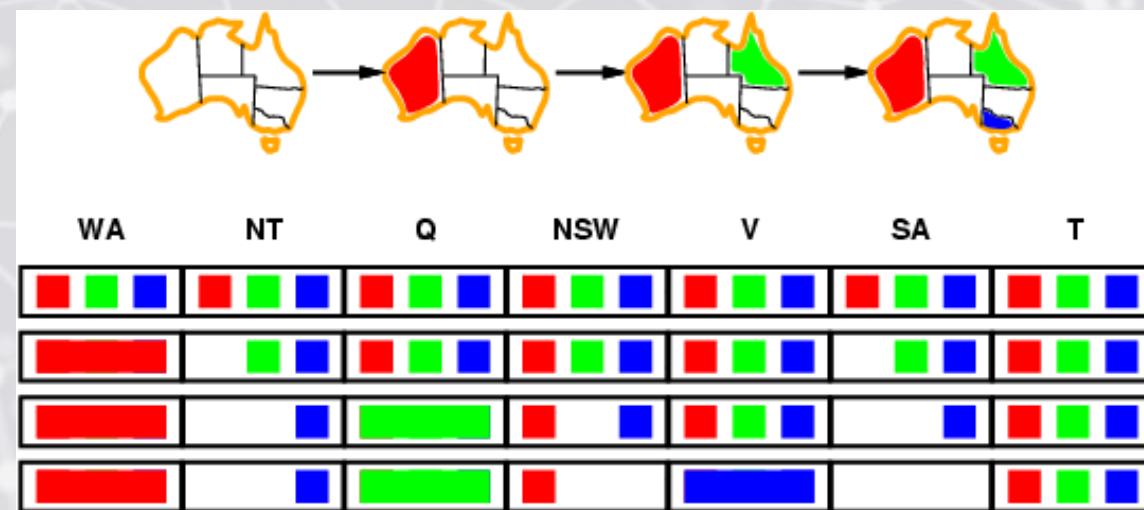
Forward checking

- Idea:
 - Keep track of remaining legal values for unassigned variables.
 - Terminate search when any variable has no legal values.



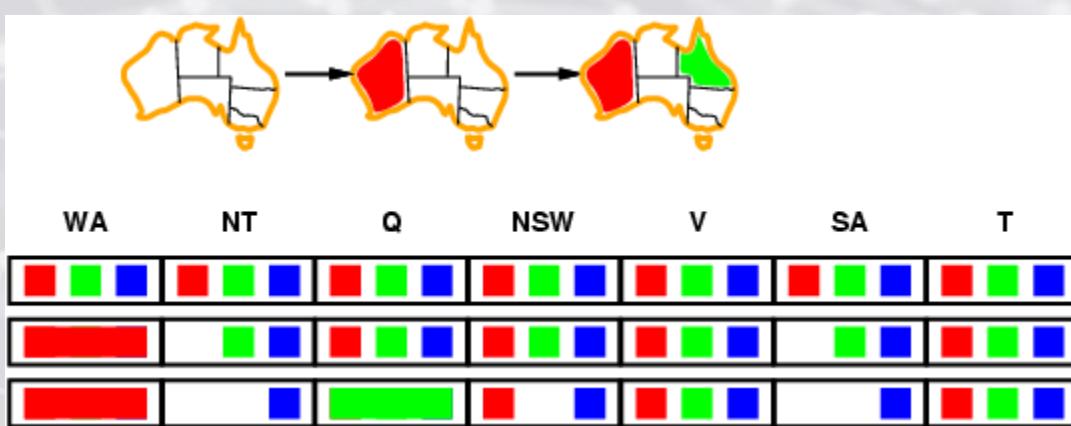
Forward checking

- Idea:
 - Keep track of remaining legal values for unassigned variables.
 - Terminate search when any variable has no legal values.



Constraint propagation

- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:



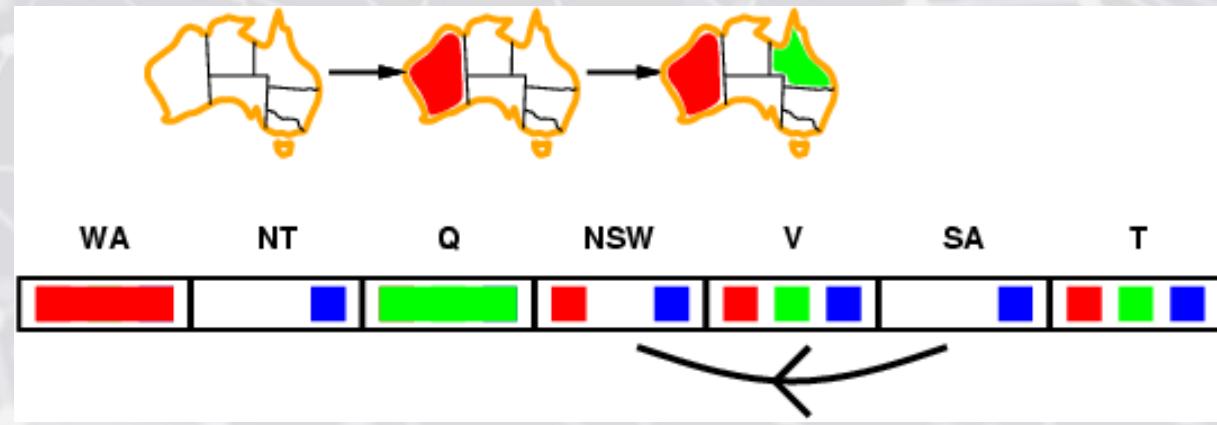
- NT and SA cannot both be blue!
 - » Constraint propagation algorithms repeatedly enforce constraints locally...

Node Consistency

- A single variable (in a CSP network) is **node-consistent** if all the values in the variable's domain satisfy the variable's unary constraints.
- E.g. Suppose South Australians dislike green; the variable SA starts with domain: {red, green, blue}, and we make it node-consistent by eliminating green, leaving SA with the reduced domain: {red, blue}.
- We say that a network is node-consistent if every variable in the network is node-consistent.

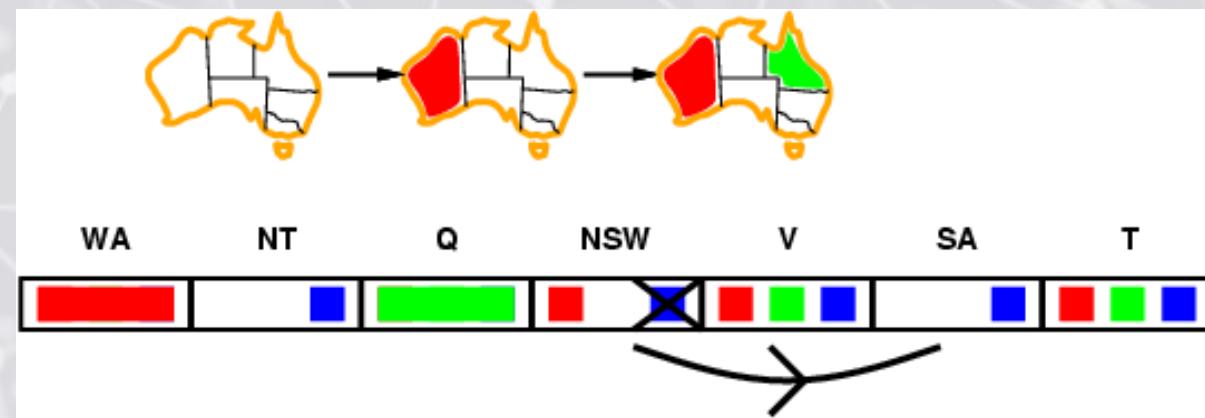
Arc consistency

- Simplest form of propagation makes each arc consistent
- $X \rightarrow Y$ is consistent iff for **every** value x of X there is **some** allowed y



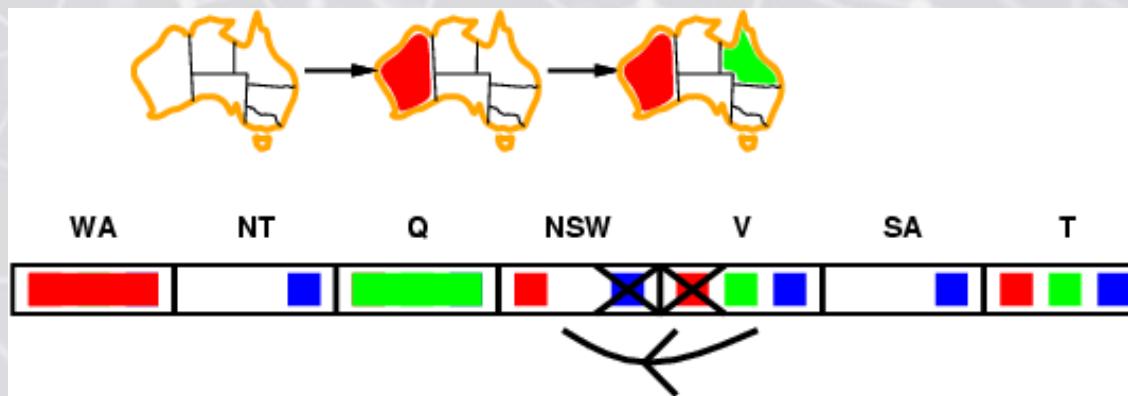
Arc consistency

- Simplest form of propagation makes each arc consistent
- $X \rightarrow Y$ is consistent iff for **every** value x of X there is **some** allowed y



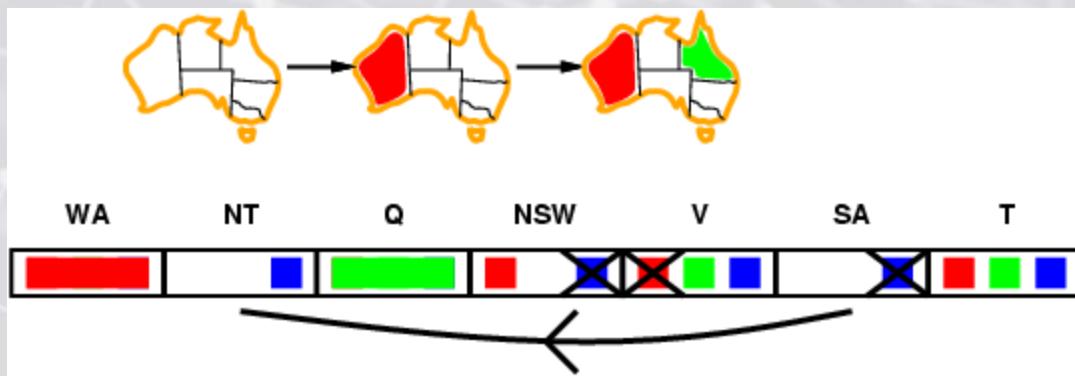
Arc consistency

- Simplest form of propagation makes each arc consistent
- $X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y



Arc consistency

- Simplest form of propagation makes each arc **consistent**
- $X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y



- If X loses a value, neighbors of X need to be rechecked
- Arc consistency detects failure earlier than forward checking

Can be run as a preprocessor or after each assignment

AC-3 Algorithm

- The most popular algorithm for arc consistency is called *AC-3*.
- The AC-3 algorithm maintains a queue of arcs to consider.
- Initially, the queue contains all the arcs in the CSP. AC-3 then pops off an arbitrary arc (X_i, X_j) from the queue and makes X_i arc-consistent with respect to X_j .

AC-3 Algorithm

- If this leaves D_i unchanged, the algorithm moves on to the next arc.
- If this revises D_i , then we add to the queue all arcs (X_k, X_i) , where X_k is a neighbor of X_i .
- We need to do this because the change in D_i might enable further reductions in the domains of D_k .
- Continue this process...
- We end up with a CSP *equivalent* to the original CSP – but the arc-consistent CSP will in most cases be much faster to search.

Arc consistency algorithm AC-3

```
function AC-3( csp) returns the CSP, possibly with reduced domains
    inputs: csp, a binary CSP with variables  $\{X_1, X_2, \dots, X_n\}$ 
    local variables: queue, a queue of arcs, initially all the arcs in csp

    while queue is not empty do
         $(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\textit{queue})$ 
        if RM-INCONSISTENT-VALUES( $X_i, X_j$ ) then
            for each  $X_k$  in NEIGHBORS[ $X_i$ ] do
                add  $(X_k, X_i)$  to queue



---


function RM-INCONSISTENT-VALUES(  $X_i, X_j$ ) returns true iff remove a value
     $removed \leftarrow false$ 
    for each  $x$  in DOMAIN[ $X_i$ ] do
        if no value  $y$  in DOMAIN[ $X_j$ ] allows  $(x,y)$  to satisfy constraint( $X_i, X_j$ )
            then delete  $x$  from DOMAIN[ $X_i$ ];  $removed \leftarrow true$ 
    return removed
```

- Time complexity: $O(\# \text{constraints} \cdot |\text{domain}|^3)$

Checking consistency of an arc is $O(|\text{domain}|^2)$

Sudoku with AC-3

- Consider the **AC-3** and **backtracking** (with **MRV – min remaining values** -- heuristic) algorithms for solving *Sudoku puzzles*.
- The objective of the game is just to fill a 9×9 grid with numerical digits so that each column, each row, and each of the nine 3×3 sub-grids (also called boxes) contains one of all of the digits 1 through 9.

	1	2	3	4	5	6	7	8	9
A		3		2		6			
B	9		3	5				1	
C		1	8	6	4				
D		8	1	2	9				
E	7							8	
F		6	7	8	2				
G		2	6	9	5				
H	8		2	3			9		
I		5	1	3					

	1	2	3	4	5	6	7	8	9
A	4	8	3	9	2	1	6	5	7
B	9	6	7	3	4	5	8	2	1
C	2	5	1	8	7	6	4	9	3
D	5	4	8	1	3	2	9	7	6
E	7	2	9	5	6	4	1	3	8
F	1	3	6	7	9	8	2	4	5
G	3	7	2	6	8	9	5	1	4
H	8	1	4	2	5	3	7	6	9
I	6	9	5	4	1	7	3	8	2

Sudoku with AC-3

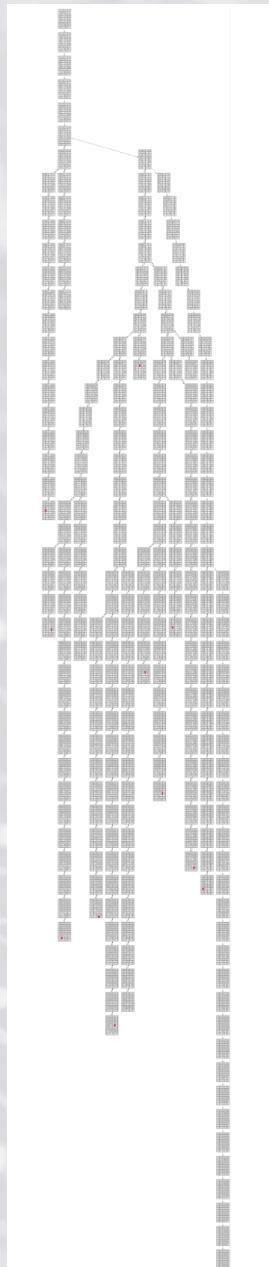
- First, the AC-3 (**arc-consistency** checking) algorithm is implemented. This algorithm will propagate the constraints and reduce the domain size of the variables by ensuring all possible (future) assignments consistent.
- Very few of the puzzles can be solved by using this algorithm alone.

Sudoku with AC-3

- First, the AC-3 (**arc-consistency** checking) algorithm is implemented. This algorithm will propagate the constraints and reduce the domain size of the variables by ensuring all possible (future) assignments consistent.
- Very few of the puzzles can be solved by using this algorithm alone.
- Next, we implement a backtracking algorithm using the **minimum remaining value** (MRV) heuristic. The order of values to be attempted for each variable can be arbitrary. When a variable is assigned, we apply **forward checking** to further reduce variables domains.

Sudoku with AC-3

			1			7		2
	3		9	5				
		1			2			3
5	9					3		1
	2						7	
7	3						9	8
8			2			1		
				8	5		6	
6		5			9			



Path Consistency

- A two-variable set $\{X_i, X_j\}$ is **path-consistent** with respect to a third variable X_m , if, for every assignment $\{X_i=a, X_j=b\}$ consistent with the constraints on $\{X_i, X_j\}$, there is an assignment to X_m that satisfies the constraints on $\{X_i, X_m\}$ and $\{X_m, X_j\}$.
- This is called path-consistency, because one can think of it as looking at a path from X_i to X_j with X_m in the middle.

k -consistency

- A CSP is k -consistent if, for any set of $k-1$ variables, and for any consistent assignment to those variables, a consistent value can always be assigned to any k th variable
- 1-consistency is node consistency.
- 2-consistency is arc consistency.
- For binary constraint networks, 3-consistency is the same as *path consistency*.
- Getting k -consistency requires time and space exponential in k .
- *Strong k -consistency* means k' -consistency for all k' from 1 to k
 - Once strong k -consistency for $k=\#\text{variables}$ has been obtained, solution can be constructed trivially
- Tradeoff between propagation and branching.
- Practitioners usually use 2-consistency and less commonly 3-consistency.

Other techniques for CSPs

- Global constraints
 - E.g., Alldiff
 - E.g., Atmost(10,P1,P2,P3), i.e., sum of the 3 vars ≤ 10
 - Special propagation algorithms
 - Bounds propagation
 - E.g., number of people on two flight D1 = [0, 165] and D2 = [0, 385]
 - Constraint that the total number of people has to be at least 420
 - Propagating bounds constraints yields D1 = [35, 165] and D2 = [255, 385]
 - ...
- Symmetry breaking (e.g. reduce search by imposing arbitrary ordering constraint).

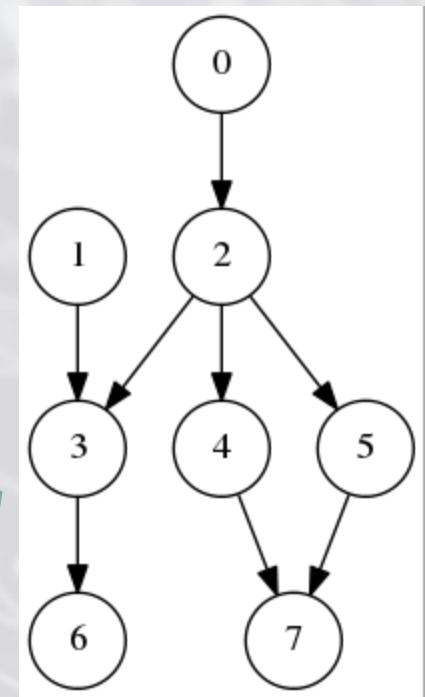
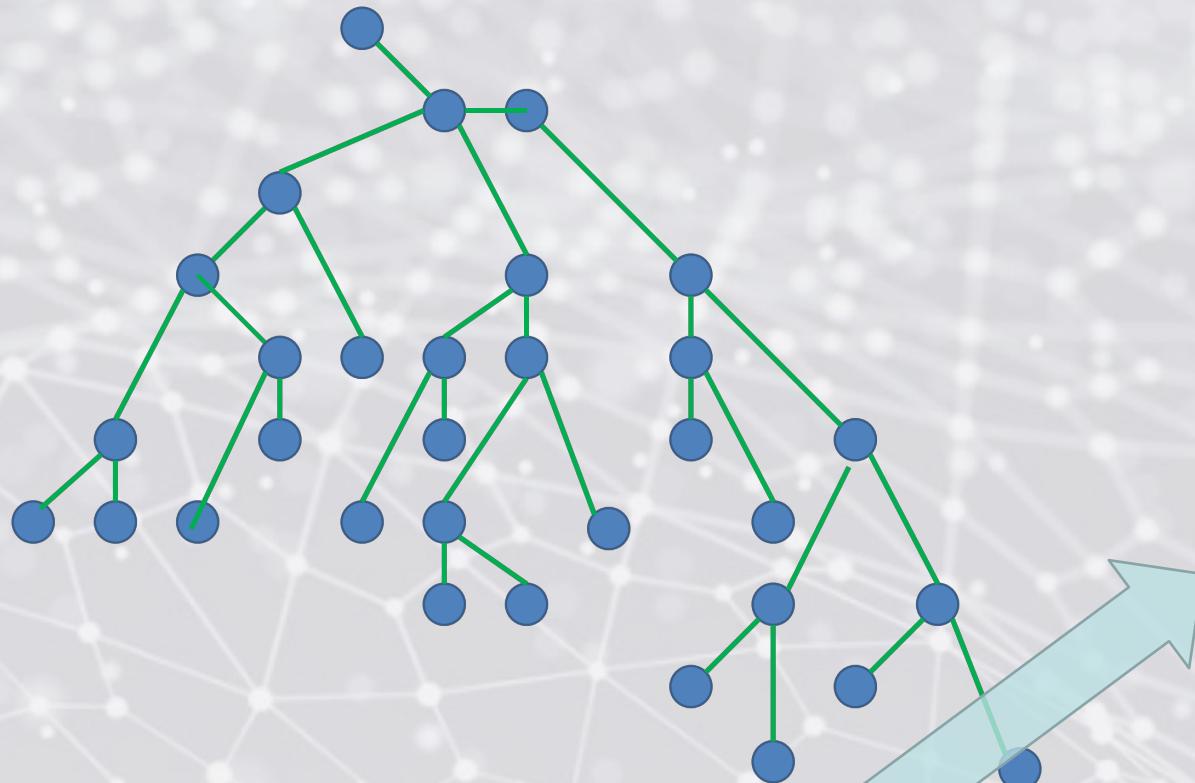
Structured CSPs

Tree-structured CSPs



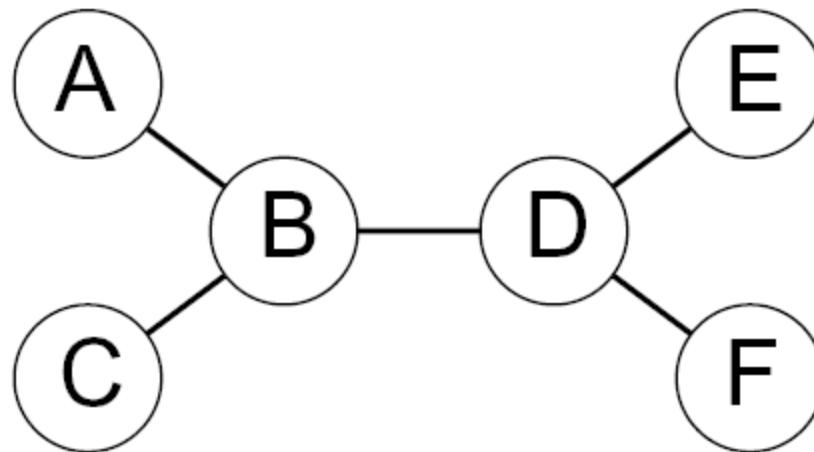
- Recall from our previous discussions about trees:
 - (*) Trees are **acyclic and connected** (i.e. they are minimally connected structures); a unique path connects any two pair of nodes; any node can serve as a “root”; trees admit of a **topological ordering**.

Tree-structured CSPs



- Recall from our previous discussions about trees:
 - (*) Trees are **acyclic and connected** (i.e. they are minimally connected structures); a unique path connects any two pair of nodes; any node can serve as a “root”; trees admit of a **topological ordering**.

Tree-structured CSPs



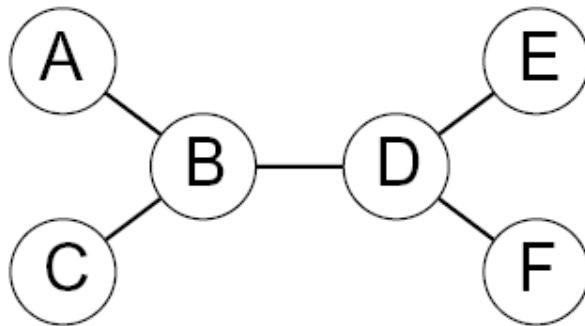
Theorem: if the constraint graph has no loops, the CSP can be solved in $O(nd^2)$ time

Compare to general CSPs, where worst-case time is $O(d^n)$

This property also applies to logical and probabilistic reasoning:
an important example of the relation between syntactic restrictions
and the complexity of reasoning.

Algorithm for tree-structured CSPs

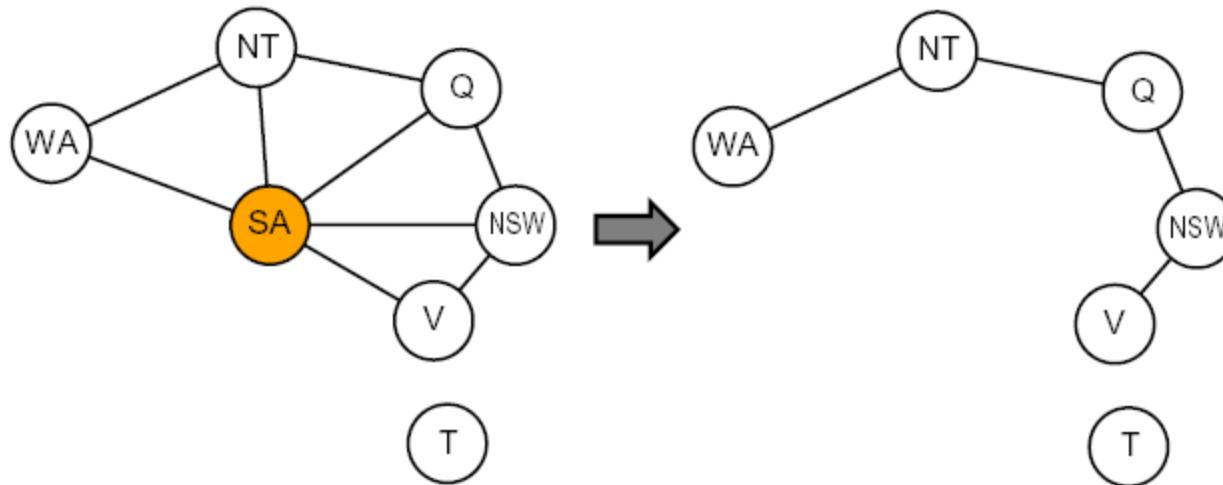
1. Choose a variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering



2. For j from n down to 2, apply $\text{REMOVEINCONSISTENT}(\text{Parent}(X_j), X_j)$
3. For j from 1 to n , assign X_j consistently with $\text{Parent}(X_j)$

Nearly tree-structured CSPs

Conditioning: instantiate a variable, prune its neighbors' domains



Cutset conditioning: instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree (Finding the minimum cutset is NP-complete.)

Cutset size c \Rightarrow runtime $O(d^c \cdot (n - c)d^2)$, very fast for small c